Tree Borrows

Neven Villani,[1] Johannes Hostert,[2] Derek Dreyer,[3] Ralf Jung[2]

PLDI'25

2025-06-19

[1]Univ. Grenoble Alpes, Verimag
[2]ETH Zurich
[3]MPI-SWS

Rust's type system enables powerful optimizations

```
fn write_both(x: &mut i32, y: &mut i32) -> i32 {

  *x = 13;

  *y = 20;

  *x

}
```

Rust's type system enables powerful optimizations

```
fn write_both(x: &mut i32, y: &mut i32) -> i32 {
    *x = 13;
    *y = 20;
    *x
}
```

mutable thus disjoint

*x is unchanged

Rust's type system enables powerful optimizations

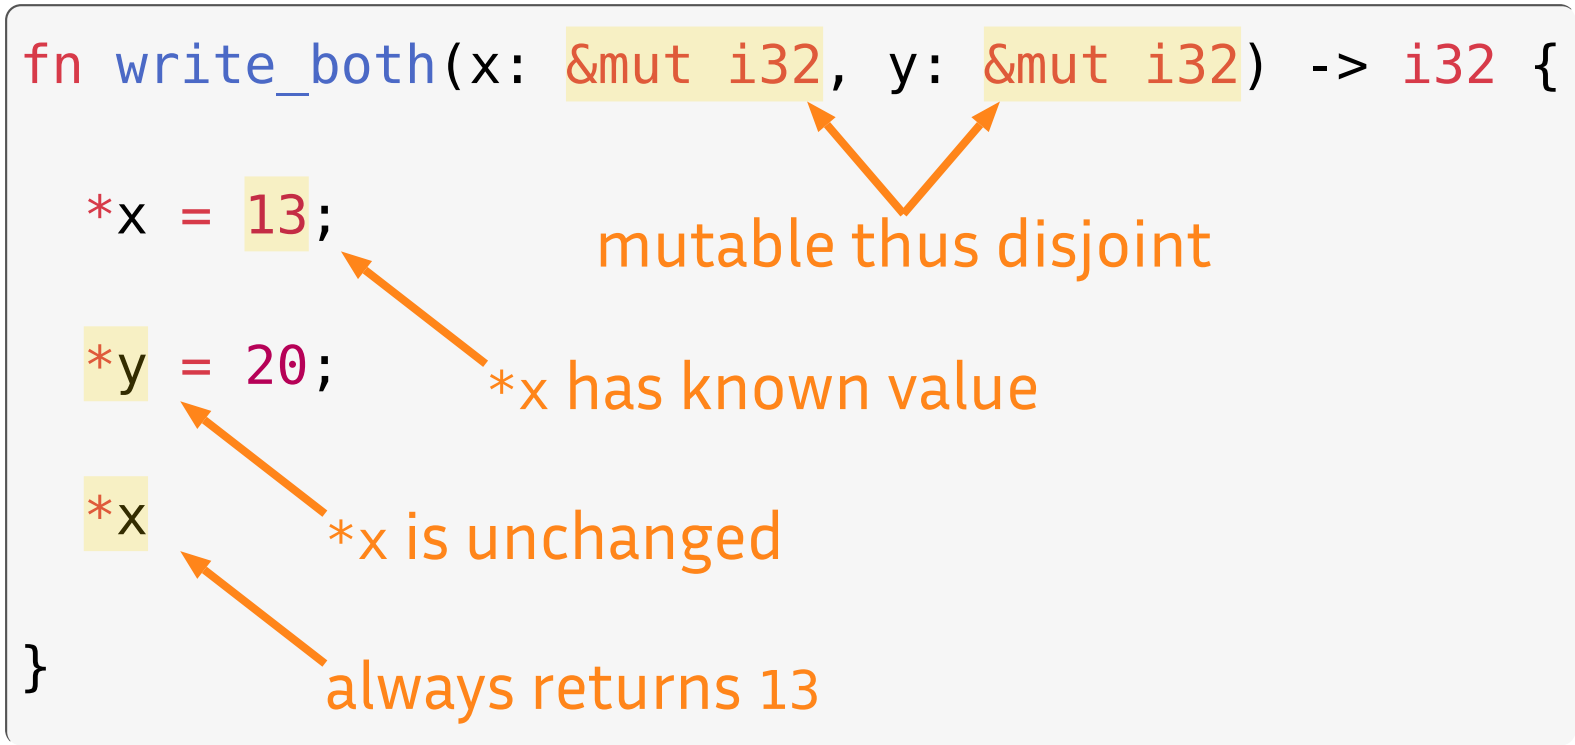

Rust's type system enables powerful optimizations

```rust
fn write_both(x: &mut i32, y: &mut i32) -> i32 {
  *x = 13;
  *y = 20;
  13 // formerly *x: one fewer load from memory
}
```

Type-level guarantees for references

`&mut` → mutation, no aliasing

`&` → aliasing, no mutation

Escape hatch: `unsafe`

Can use **unchecked operations** to do **low-level manipulations**

```
unsafe {
  // Code within this block can effectively
  // bypass some parts of the typechecker.
  ...
}
```

Within `unsafe` it is **the programmer's responsibility** to check

- that pointers are non-null
- that memory is initialized
- absence of data races
- ...

} violations trigger UB (Undefined Behavior)

What if `unsafe` code is misused ?

```rust
fn write_both(x: &mut i32, y: &mut i32) -> i32 {
  *x = 13;
  *y = 20;
  *x
}

fn main() {
  let mut root = 42;
  let ptr = &raw mut root;
  let x = unsafe { &mut *ptr };
  let y = unsafe { &mut *ptr };
  println!("{}", write_both(x, y)); // prints 20
}
```

What if unsafe code is misused ?

```rust
fn write_both(x: &mut i32, y: &mut i32) -> i32 {
  *x = 13;
  *y = 20;
  *x
}

fn main() {
  let mut root = 42;
  let ptr = &raw mut root;
  let x = unsafe { &mut *ptr };
  let y = unsafe { &mut *ptr };
  println!("{}", write_both(x, y)); // prints 20
}
```

What if unsafe code is misused ?

```
fn write_both(x: &mut i32, y: &mut i32) -> i32 {
  *x = 13;
  *y = 20;
  *x
}

fn main() {
  let mut root = 42;
  let ptr = &raw mut root;
  let x = unsafe { &mut *ptr };
  let y = unsafe { &mut *ptr };
  println!("{}", write_both(x, y)); // prints 20
}
```

What if unsafe code is misused ?

```
fn write_both(x: &mut i32, y: &mut i32) -> i32 {
  *x = 13;
  *y = 20;
  *x
}

fn main() {
  let mut root = 42;
  let ptr = &raw mut root;
  let x = unsafe { &mut *ptr };
  let y = unsafe { &mut *ptr };
  println!("{}", write_both(x, y)); // prints 20
}
```

What if `unsafe` code is misused ?

```
fn write_both(x: &mut i32, y: &mut i32) -> i32 {
  *x = 13;
  *y = 20;
  13
}

fn main() {
  let mut root = 42;
  let ptr = &raw mut root;
  let x = unsafe { &mut *ptr };
  let y = unsafe { &mut *ptr };
  println!("{}", write_both(x, y)); // prints 13
}
```

What if `unsafe` code is misused ?

```
fn write_both(x: &mut i32, y: &mut i32) -> i32 {
  *x = 13;
  *y = 20;
  *x
}
```

`unsafe` code can break the assumptions that optimizations need!

```
  let ptr = &raw mut root;
  let x = unsafe { &mut *ptr };
  let y = unsafe { &mut *ptr };
  println!("{}", write_both(x, y)); // prints 20
}
```

Expanding our notion of UB

Within `unsafe` it is **the programmer's responsibility** to check

- that pointers are non-null
- that memory is initialized
- absence of data races
- compliance with aliasing rules *NEW!*

} violations trigger UB

Expanding our notion of UB

Within `unsafe` it is **the programmer's responsibility** to check

- that pointers are non-null
- that memory is initialized
- absence of data races
- compliance with aliasing rules *NEW!*

} violations trigger UB

Tree Borrows (TB): defines those aliasing rules

Compiler **assumes absence of UB**, exploits this for optimizations

Expanding our notion of UB

Within `unsafe` it is **the programmer's responsibility** to check

- that pointers are non-null
- that memory is initialized

Sounds familiar?

TB is the successor of **Stacked Borrows**, which has the same purpose.

Stacked Borrows (SB)

[Jung et al., POPL'20]

In safe Rust, the Borrow Checker makes borrows well-bracketed.
Stacked Borrows extends the well-bracketedness to `unsafe`.

```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
```

Desired outcome: UB

In safe Rust, the Borrow Checker makes borrows well-bracketed.
Stacked Borrows extends the well-bracketedness to `unsafe`.

```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
```

- new stack at `root`

Desired outcome: UB

In safe Rust, the Borrow Checker makes borrows well-bracketed.
Stacked Borrows extends the well-bracketedness to `unsafe`.

```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
```

- ✓ `root` is at the top

Desired outcome: UB

In safe Rust, the Borrow Checker makes borrows well-bracketed.
Stacked Borrows extends the well-bracketedness to `unsafe`.

```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
```

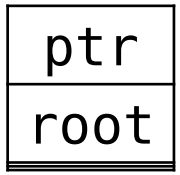

- ✓ `root` is at the top
- push `ptr`

Desired outcome: UB

In safe Rust, the Borrow Checker makes borrows well-bracketed.
Stacked Borrows extends the well-bracketedness to `unsafe`.

```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
```

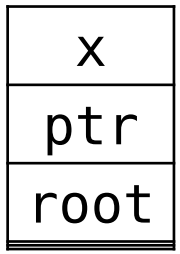

- ✓ `ptr` is at the top
- push `x`

Desired outcome: UB

In safe Rust, the Borrow Checker makes borrows well-bracketed.
Stacked Borrows extends the well-bracketedness to `unsafe`.

```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
```

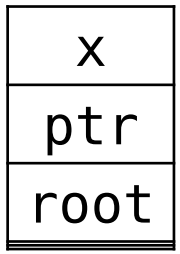

- pop until `ptr` is at the top

Desired outcome: UB

In safe Rust, the Borrow Checker makes borrows well-bracketed.
Stacked Borrows extends the well-bracketedness to `unsafe`.

```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
```

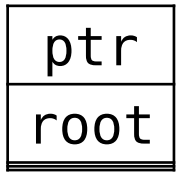

- pop until `ptr` is at the top

Desired outcome: UB

In safe Rust, the Borrow Checker makes borrows well-bracketed.
Stacked Borrows extends the well-bracketedness to `unsafe`.

```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
```

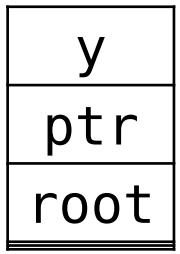

- pop until `ptr` is at the top
- push `y`

Desired outcome: UB

In safe Rust, the Borrow Checker makes borrows well-bracketed.
Stacked Borrows extends the well-bracketedness to `unsafe`.

```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
```

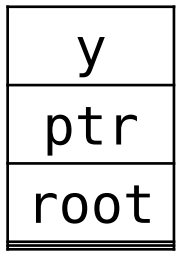

- search for x

Desired outcome: UB

In safe Rust, the Borrow Checker makes borrows well-bracketed.
Stacked Borrows extends the well-bracketedness to `unsafe`.

```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
```

UB!

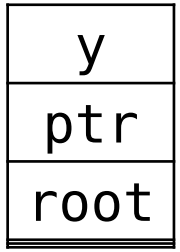

Can't use x if it is not in the stack

Desired outcome: UB

SB was **implemented** in Miri (official interpreter and UB detector)

$\rightarrow$ included in many projects' CI

$\rightarrow$ many bugs detected (e.g. in stdlib)

SB was **implemented** in Miri (official interpreter and UB detector)

→ included in many projects' CI

→ many bugs detected (e.g. in stdlib)

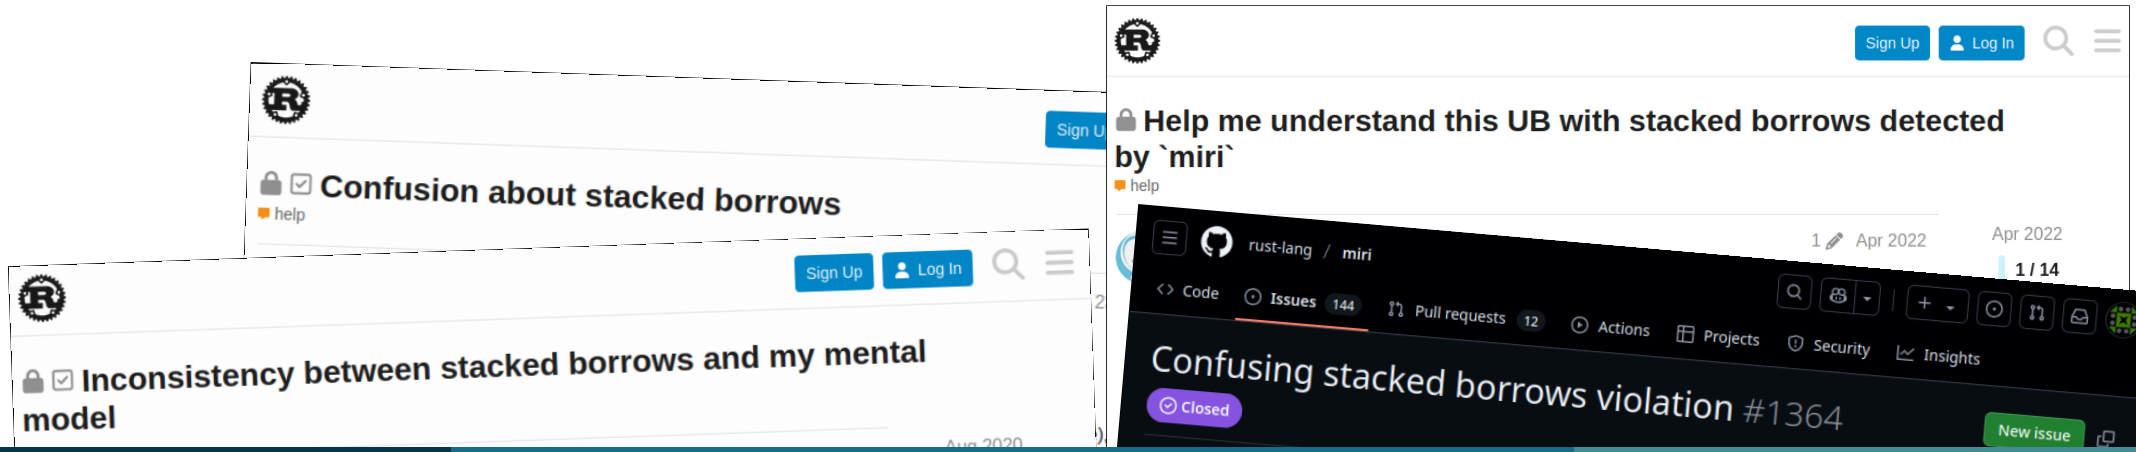

SB was **implemented** in Miri (official interpreter and UB detector)

→ included in many projects' CI

→ many bugs detected (e.g. in stdlib)

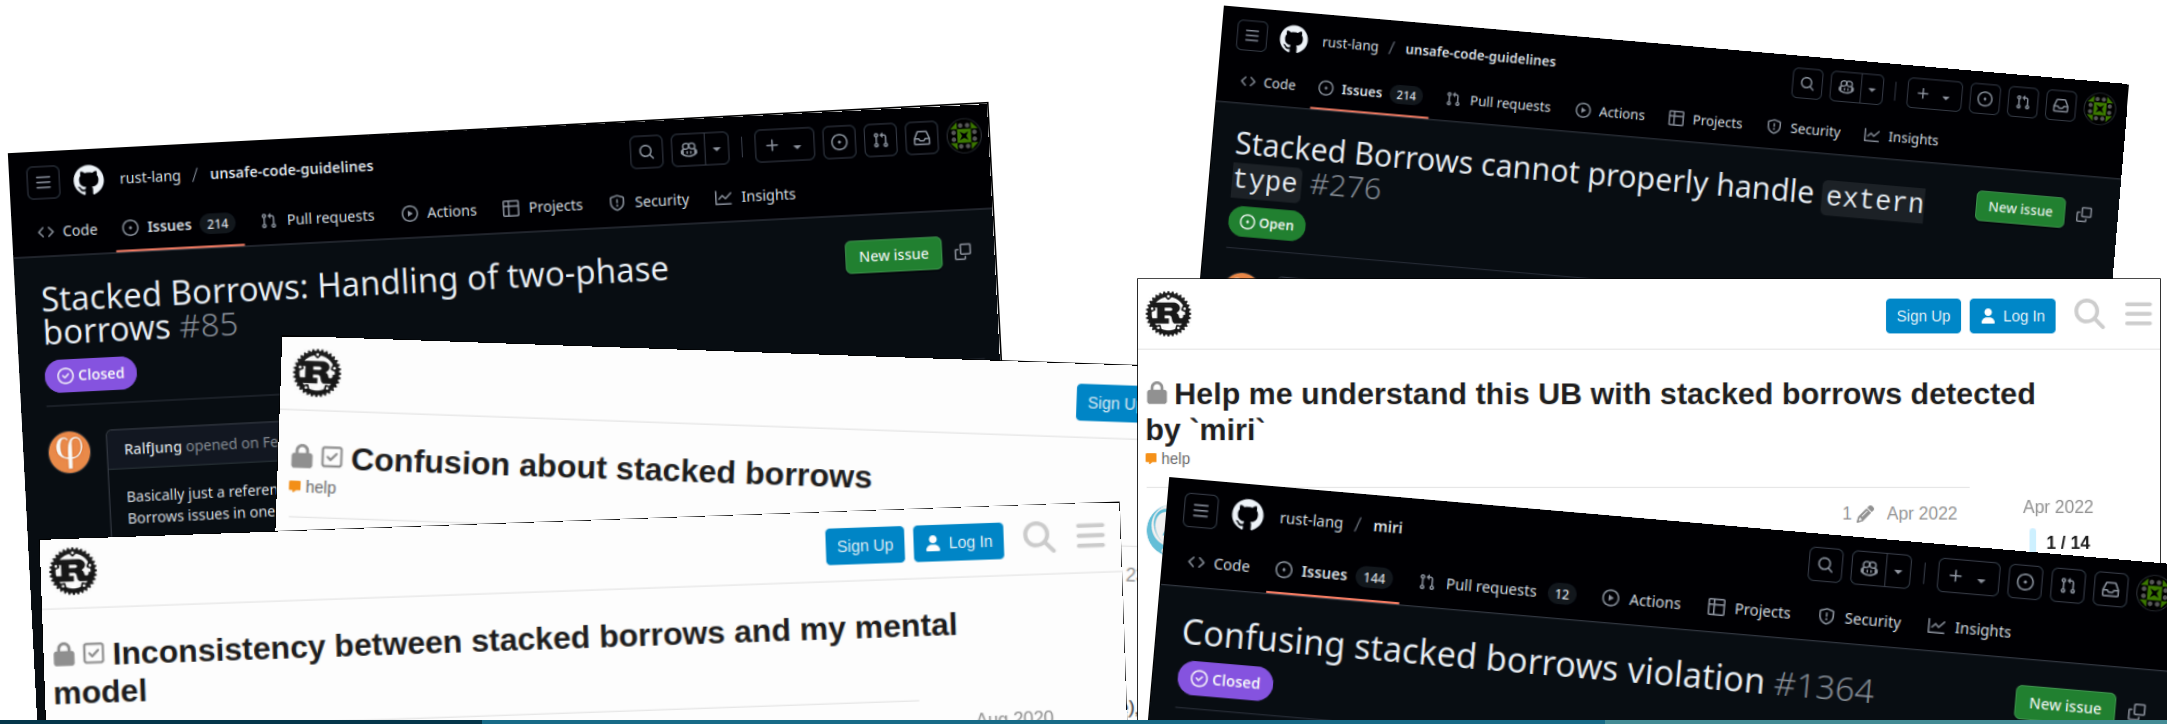

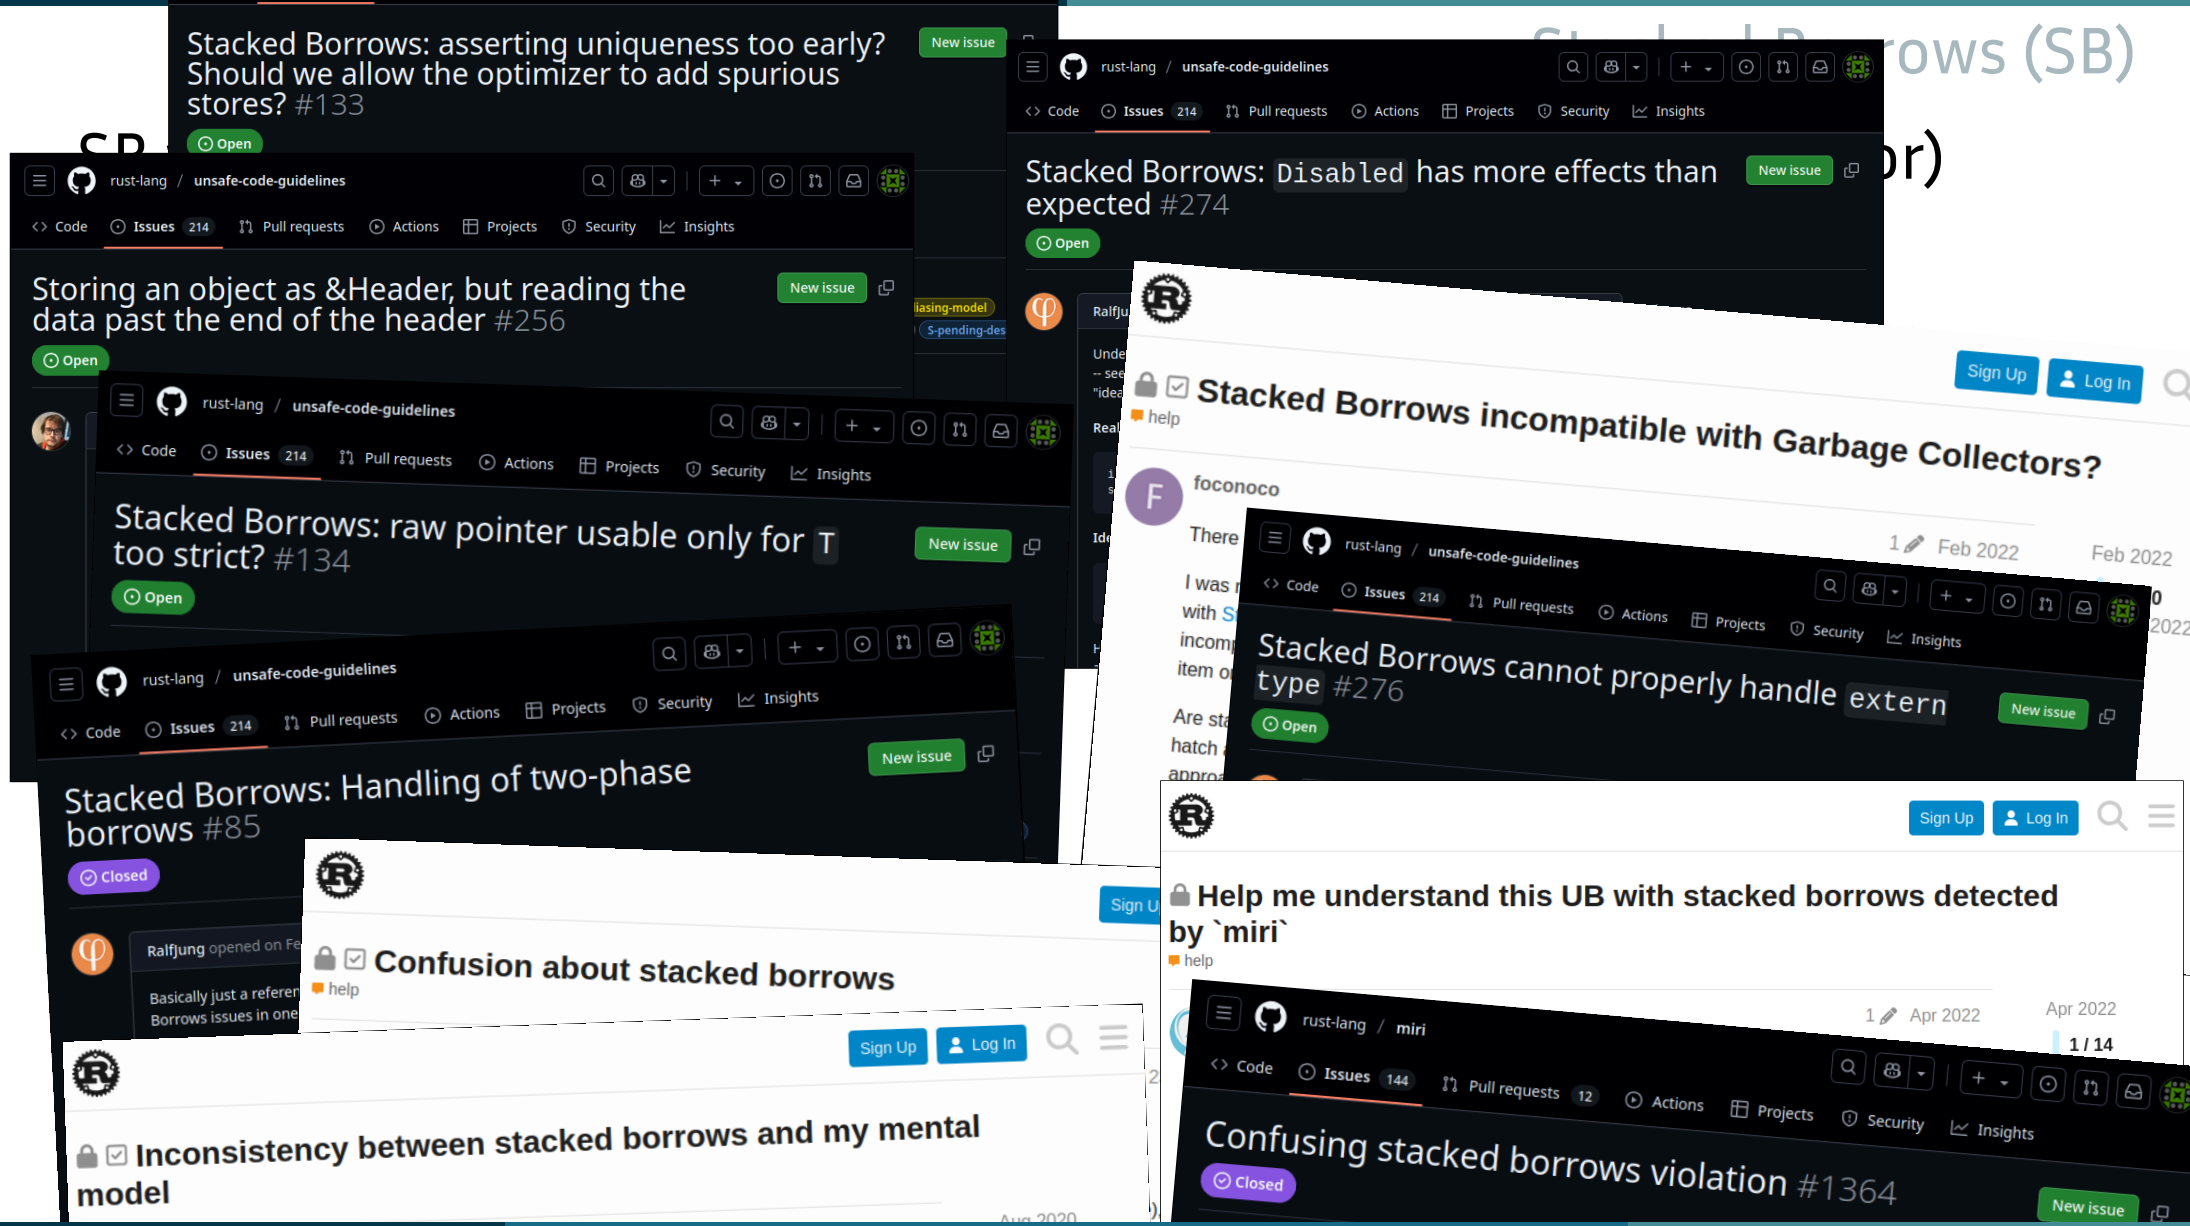
Stacked Borrows: asserting uniqueness too early? Should we allow the optimizer to add spurious stores? #133
New issue
Open
rust-lang / unsafe-code-guidelines
Code Issues 214 Pull requests Actions Projects Security Insights
Storing an object as &Header, but reading the data past the end of the header #256
New issue
Open
Stacked Borrows: raw pointer usable only for T too strict? #134
New issue
Open
Stacked Borrows: Handling of two-phase borrows #85
Closed
Ralfjung opened on
Basically just a referen
Borrows issues in one
Confusion about stacked borrows
help
Sign Up
Log In
Inconsistency between stacked borrows and my mental model
Aug 2020
Stacked Borrows: Disabled has more effects than expected #274
New issue
Open
Stacked Borrows incompatible with Garbage Collectors?
help
foconoco
There
I was
with S
incom
item o
Are sta
hatch
appro
1 Feb 2022
Feb 2022
Stacked Borrows cannot properly handle extern type #276
New issue
Open
Help me understand this UB with stacked borrows detected by `miri`
help
rust-lang / miri
Code Issues 144 Pull requests 12 Actions Projects Security Insights
1 Apr 2022
Apr 2022
1 / 14
Confusing stacked borrows violation #1364
Closed
New issue

Anecdotal evidence supported by data:

- analysis of 30 000 libraries
- 6000+ tests have aliasing UB under Stacked Borrows (leading cause of UB)

Tree Borrows allows much more code

Tree Borrows uses a **tree** instead of a stack to track borrows

Out of 30 000 most downloaded libraries,
54% fewer tests with aliasing UB when using Tree Borrows

Tree Borrows allows much more code

Tree Borrows uses a **tree** instead of a stack to track borrows

Out of 30 000 most downloaded libraries,
54% fewer tests with aliasing UB when using Tree Borrows

Fixes known technical limitations of SB,
incl. 2-phase borrows, extern types, **pointer offsets**

From Stacks to Trees

```
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];
```

```
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];
```

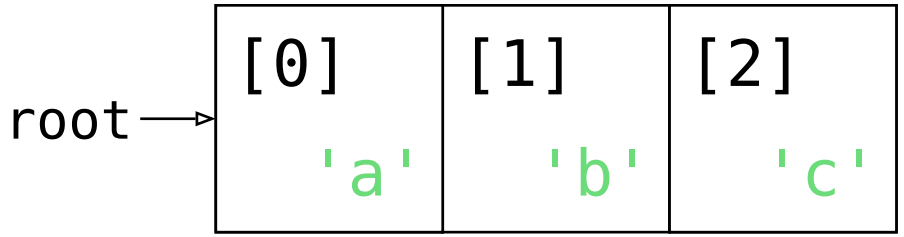

```
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];
```

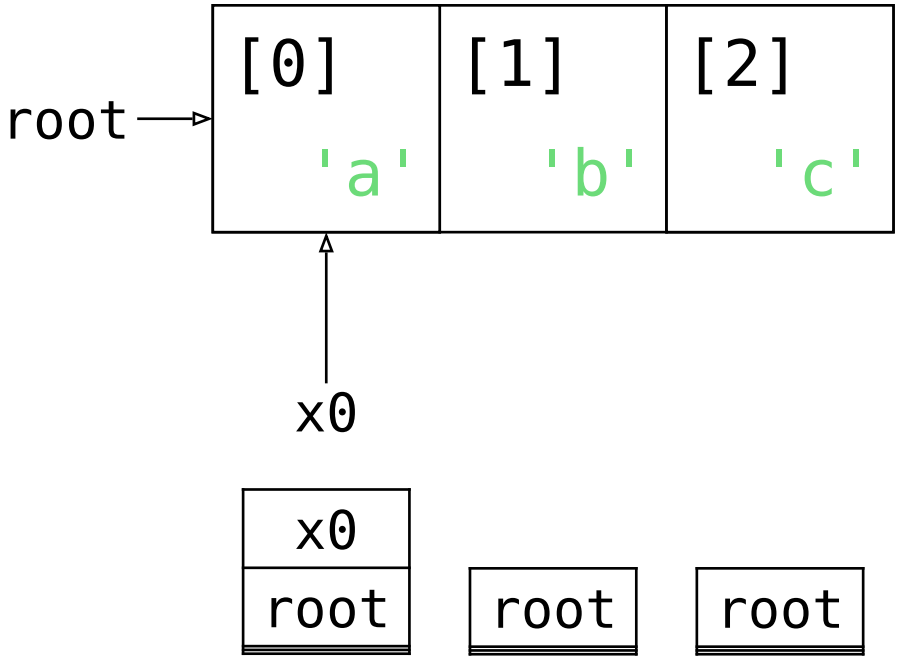

```
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];
```

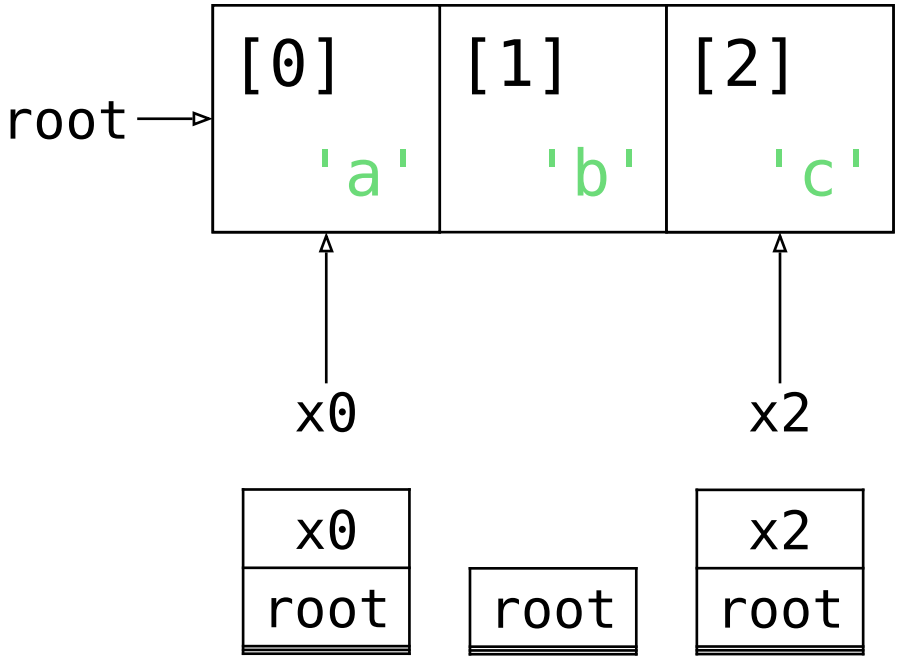

```
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];

// Scenario 1
unsafe { *x0.add(1) = 'z'; }
```

Desired outcome: not UB

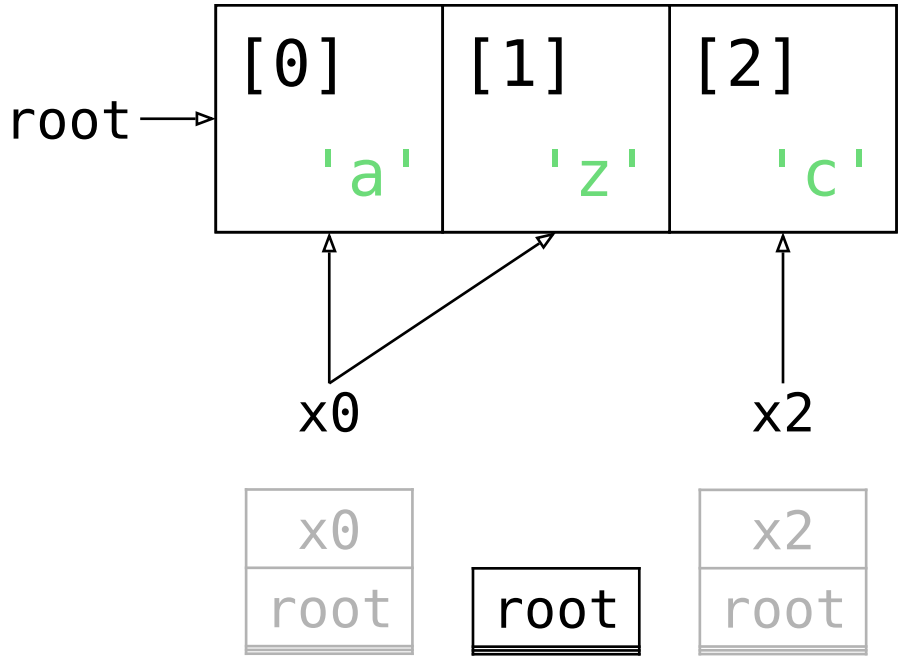

```rust
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];

// Scenario 1
unsafe { *x0.add(1) = 'z'; }
```

Desired outcome: not UB

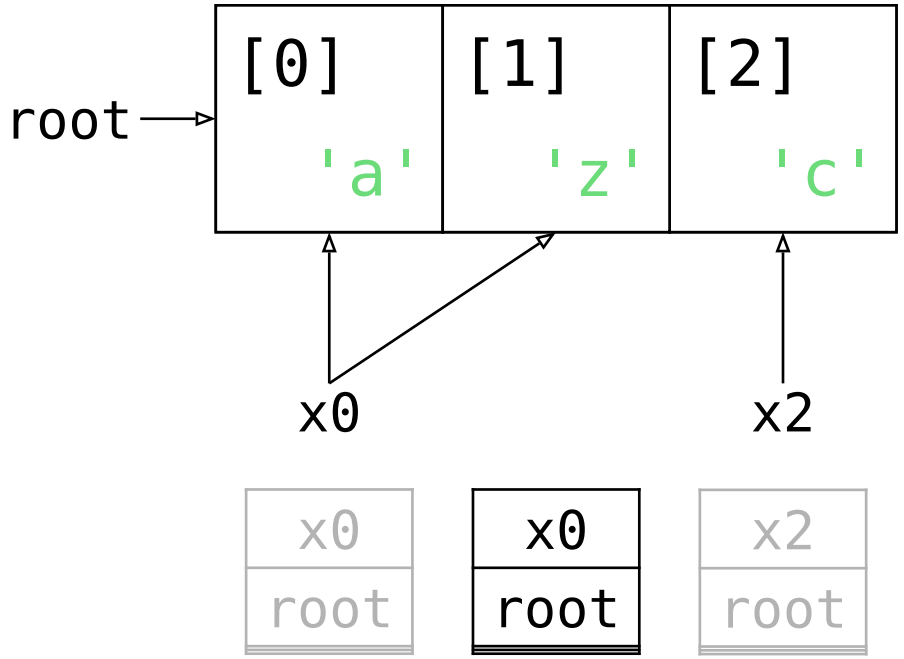

```rust
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];

// Scenario 2
unsafe { *x2.sub(1) = 'z'; }
```

Desired outcome: not UB

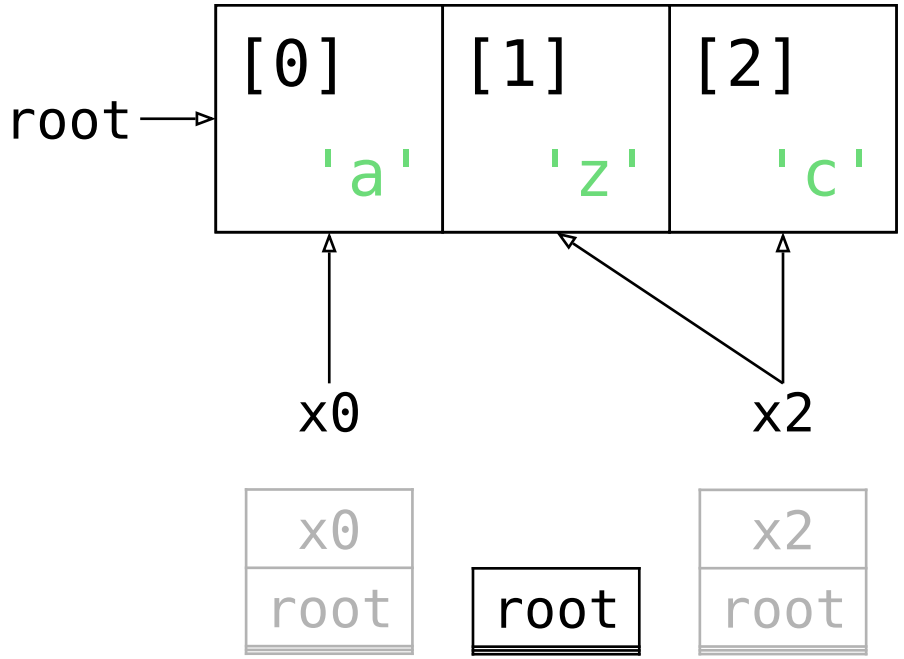

```rust
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];

// Scenario 2
unsafe { *x2.sub(1) = 'z'; }
```

Desired outcome: not UB

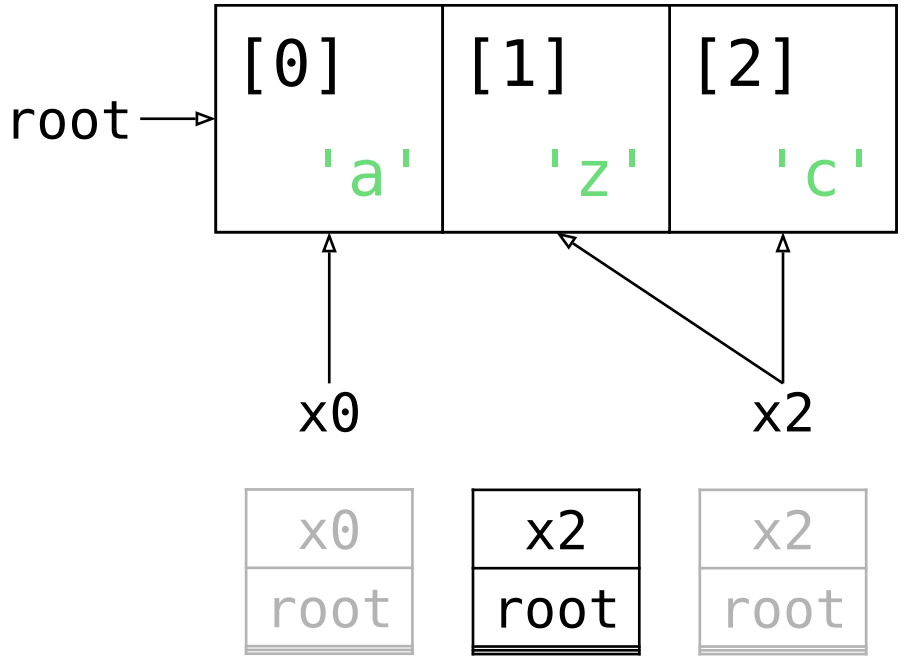

```rust
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];

// Scenario 1 or 2
unsafe { *??? = 'z'; }
```

Desired outcome: not UB

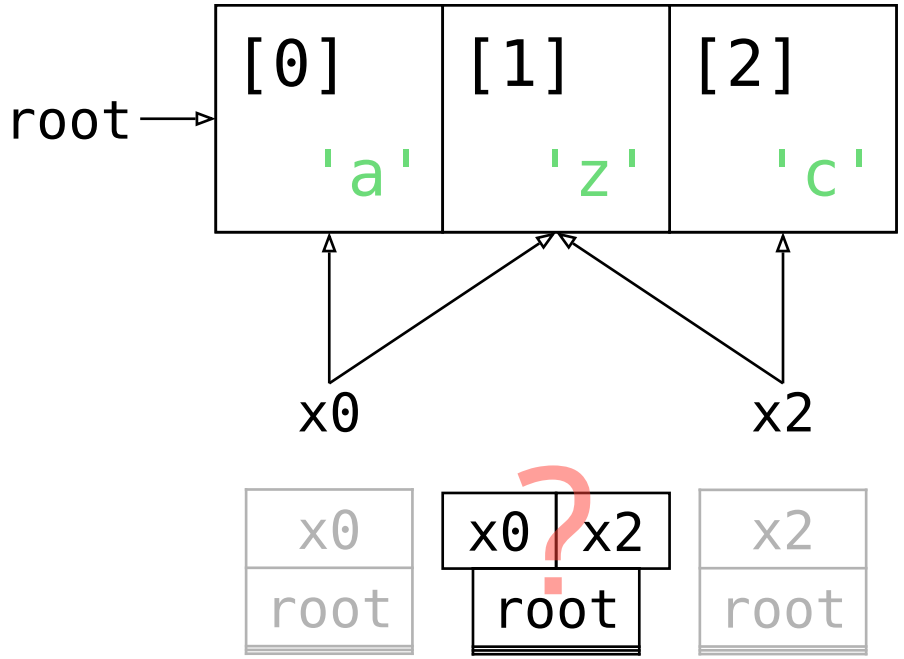

```
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];
```

Desired outcome: not UB

```rust
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];
```

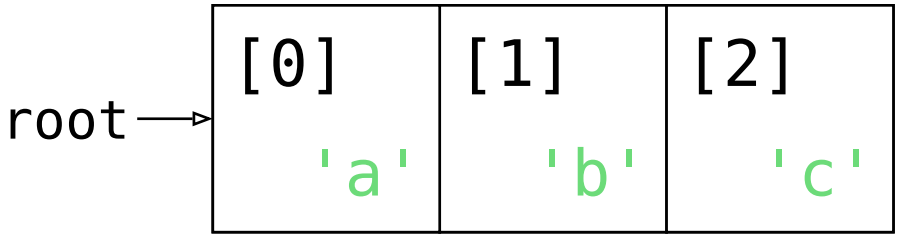

Desired outcome: not UB

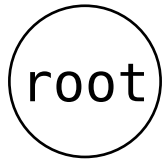

```rust
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];
```

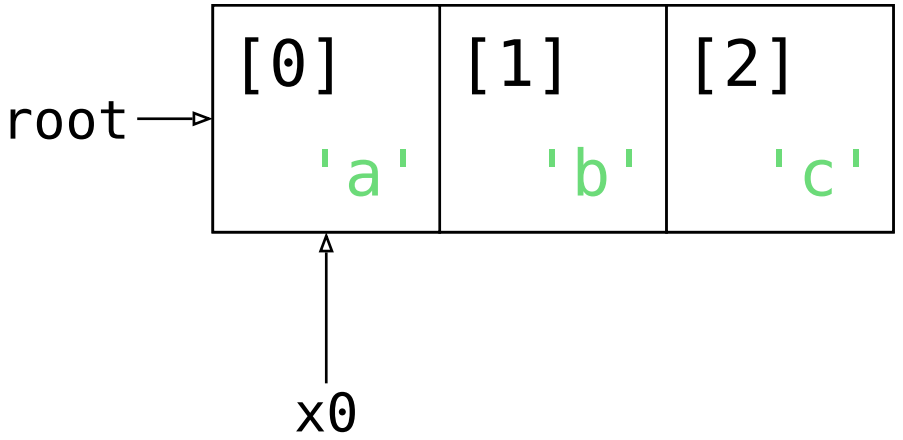

Desired outcome: not UB

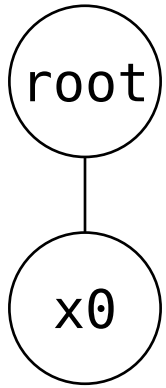

```
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];
```

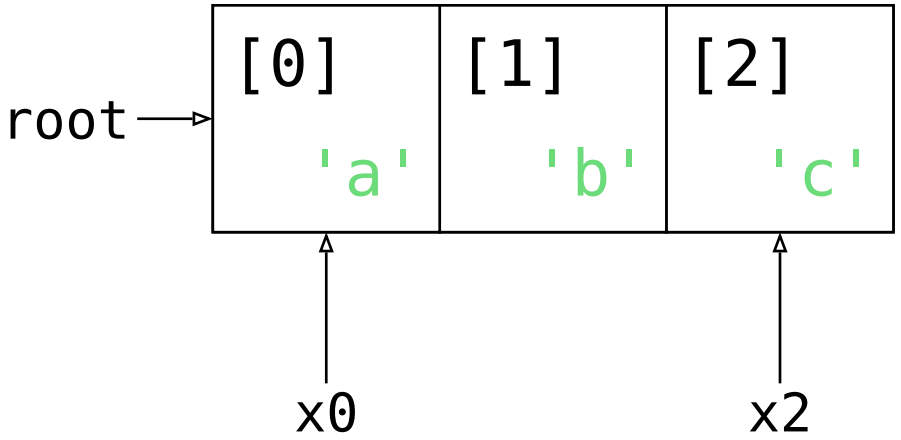

Desired outcome: not UB

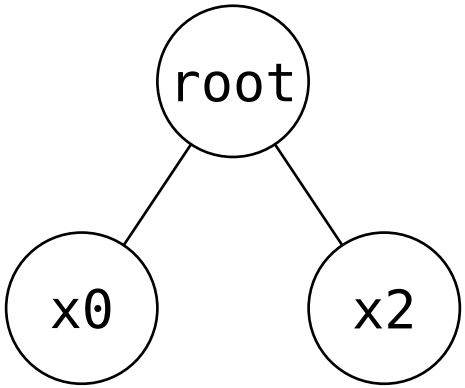

```
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];
```

Desired outcome: not UB

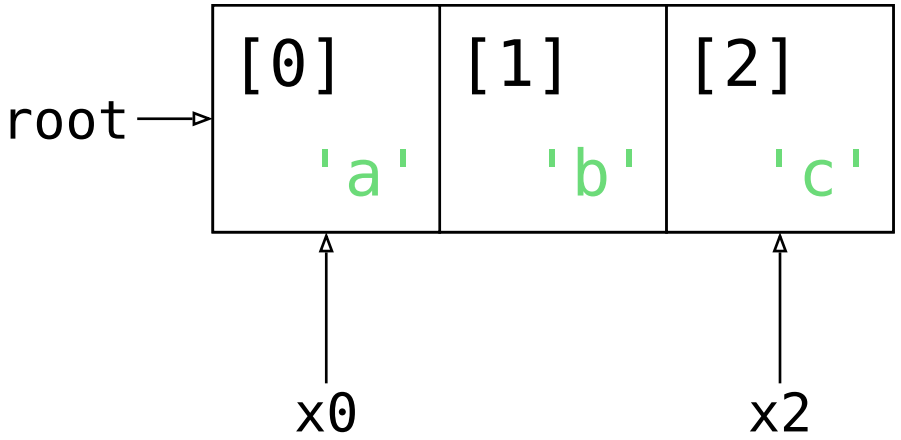

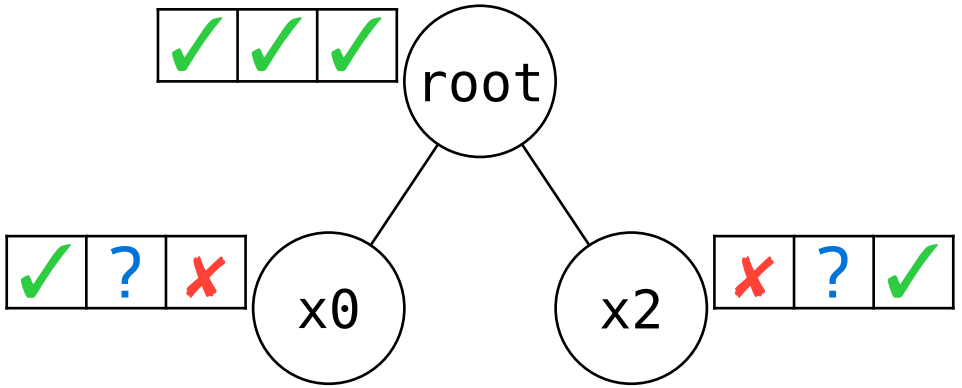

```rust
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];

// Scenario 1
unsafe { *x0.add(1) = 'z' };
```

Desired outcome: not UB

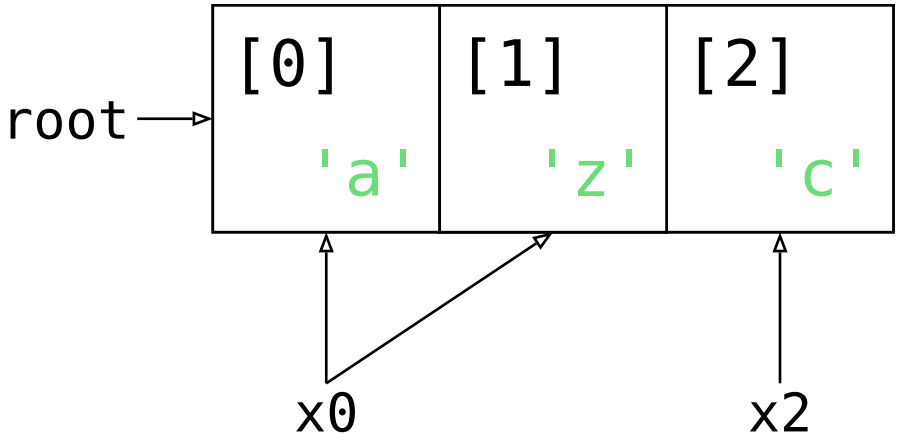

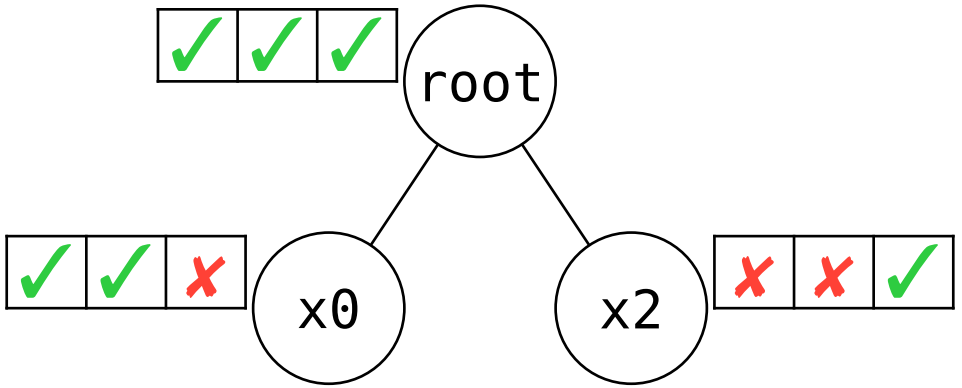

```rust
let mut root = vec!['a','b','c'];
let x0 = &raw mut root[0];
let x2 = &raw mut root[2];

// Scenario 2
unsafe { *x2.sub(1) = 'z'; }
```

Desired outcome: not UB

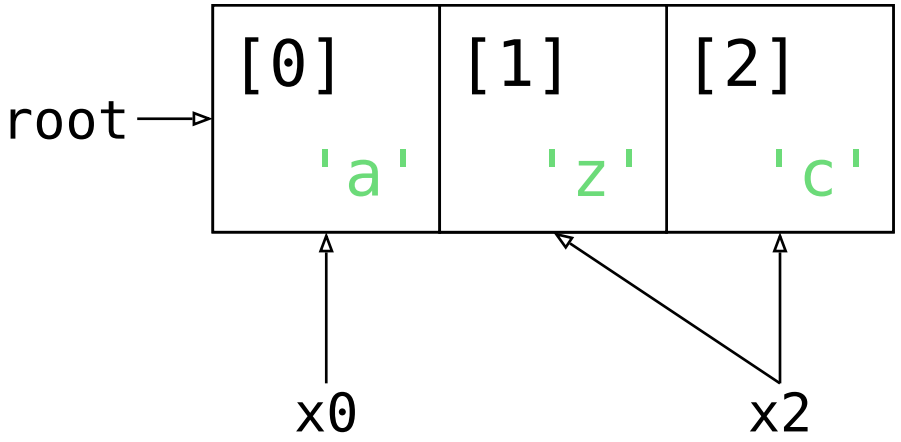

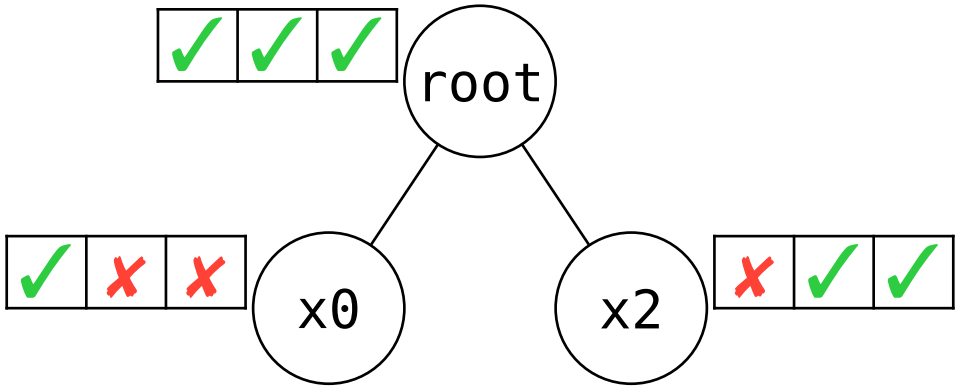

A second look at the motivating example


```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
*y = 20;
```

Desired outcome: UB

A second look at the motivating example


```rust
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
*y = 20;
```

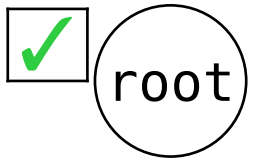

Desired outcome: UB

A second look at the motivating example


```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
*y = 20;
```

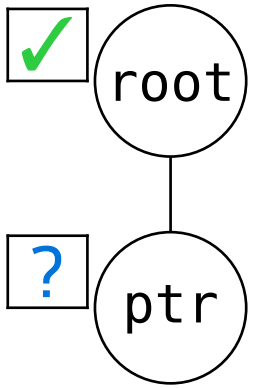

Desired outcome: UB

A second look at the motivating example


```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
*y = 20;
```

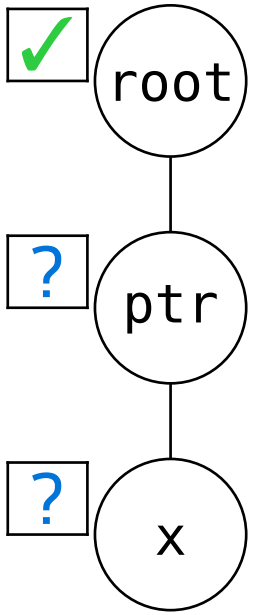

Desired outcome: UB

A second look at the motivating example


```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
*y = 20;
```

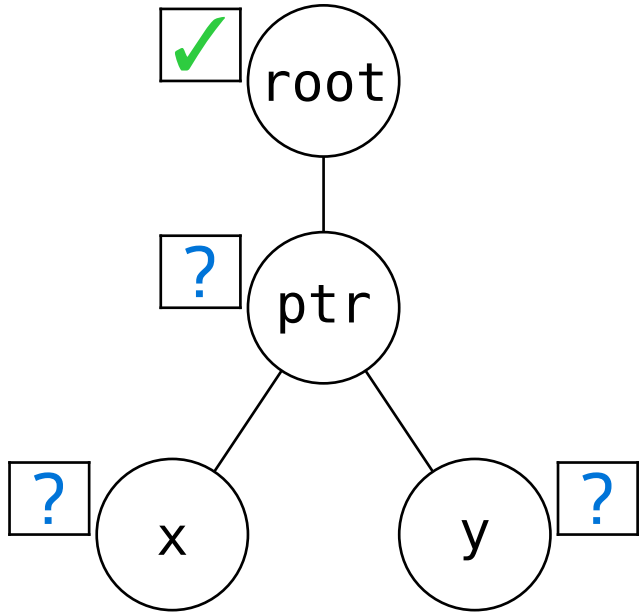

Desired outcome: UB

A second look at the motivating example


```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
*y = 20;
```

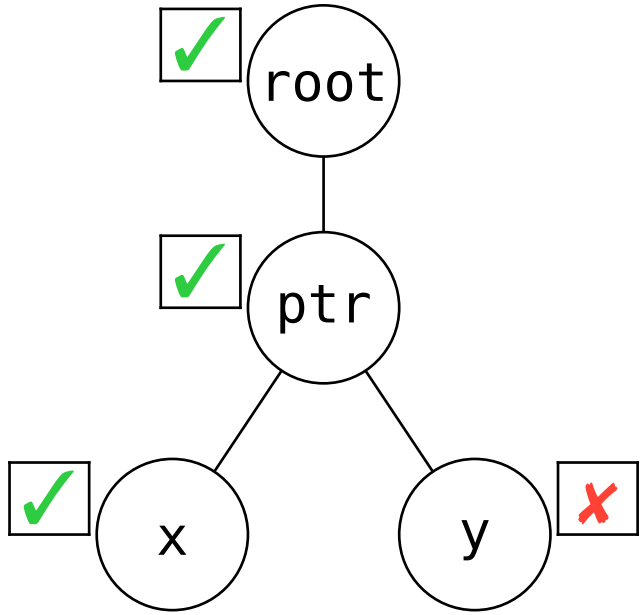

Desired outcome: UB

A second look at the motivating example


```
let mut root = 42;
let ptr = &raw mut root;
let x = unsafe { &mut *ptr };
let y = unsafe { &mut *ptr };
// inline write_both(x, y):
*x = 13;
*y = 20;
```

UB!

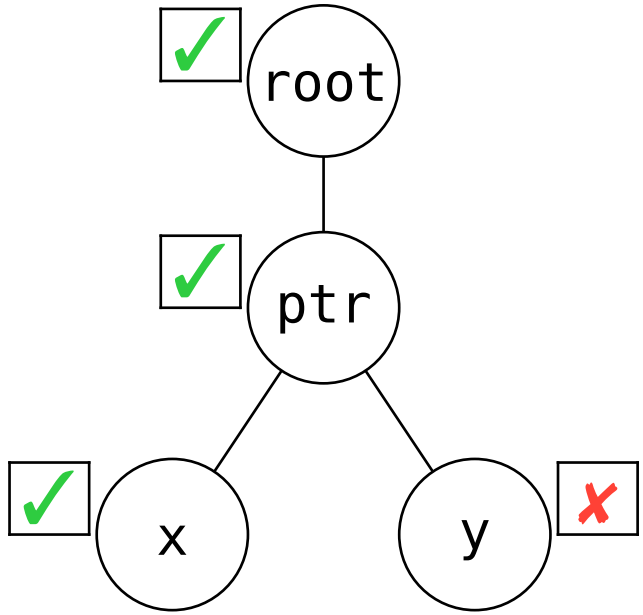

Desired outcome: UB

Evaluation

TB should enable desired optimizations

i.e. have enough UB to rule out problematic patterns

- formalized in Rocq (+Simuliris)
- a selection of optimizations proven
  - ✓ delete read through `&mut` or `&`
  - ✓ insert read through `&` in function
  - ✓ move read down for `&mut` or `&` in function

  …

It should be possible to write `unsafe` code free of UB

i.e. UB should be predictable and not too common

54% fewer tests have aliasing UB according to Tree Borrows

Only 31 ($< 0.01\%$) tests are regressions, all easily fixable.

(Out of 30 000 libraries, 400 000+ working tests)

"Tree Borrows accepts more real-world programs that call foreign functions than Stacked Borrows due to differences in handling pointer arithmetic."

A Study of Undefined Behavior Across Foreign Function Boundaries in Rust Libraries, by I. McCormack, J. Sunshine, J. Aldrich @ ICSE'25

Conclusion

Try it out: supported by Miri

Also on the Rust Playground (`play.rust-lang.org`)

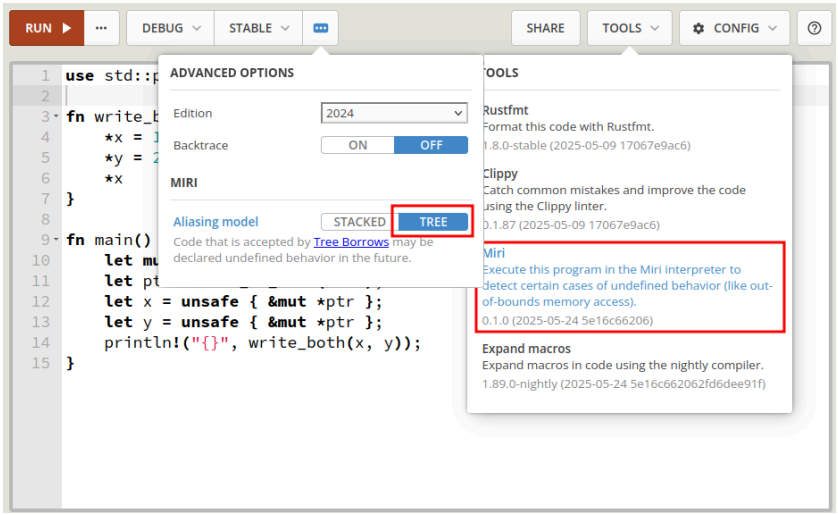

Postdoc positions available

@ ETH Zurich `ralf.jung@inf.ethz.ch`

@ MPI-SWS `dreyer@mpi-sws.org`

Learn more:

`plf.inf.ethz.ch/research/pldi25-tree-borrows.html`

Includes e.g. handling of raw pointers and interior mutability.